



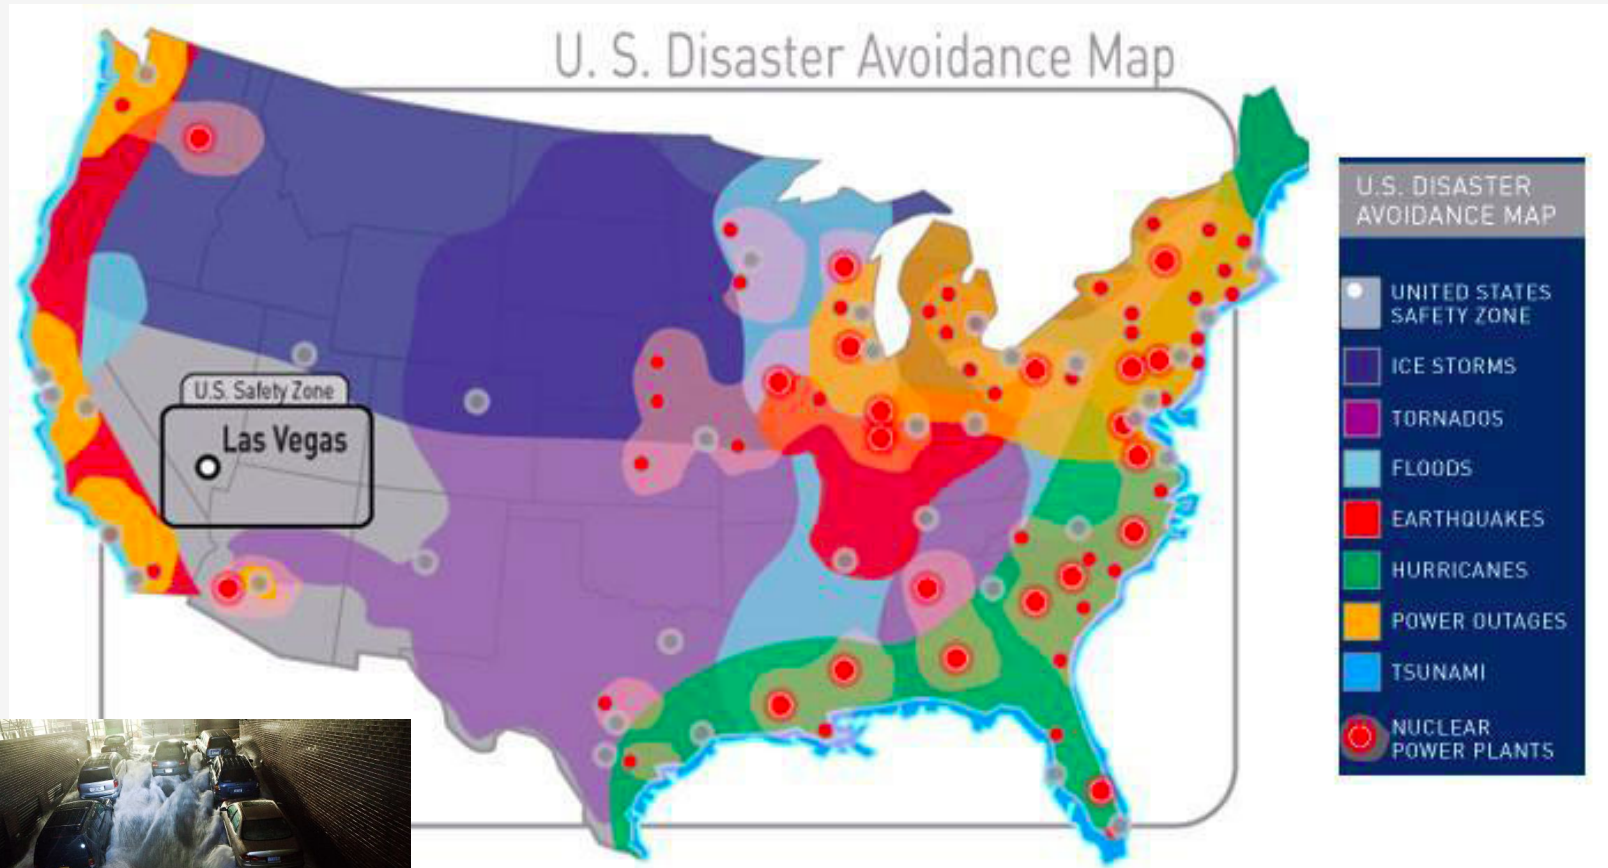
#ParisMongoDB

MongoDB Backup Overview

Sam Weaver

Solution Architect, MongoDB

Disasters do happen



Sometimes they are our fault



Availability and Durability

- Risks
 - Storage failure
 - Power outage
 - Programmer error
 - Hardware failure
 - Data centre failure
 - Cyber attack
- Storage level options
- MongoDB Backup Service

Addressing Risks

- Relative to any particular risk
 - How much data can you afford to lose? (RPO)
 - How long can you afford to be offline? (RTO)
 - What price are you willing to pay to remove the risk?
- Solutions
 - Replication
 - Application engineering?
 - Backup

Replication

- Built into MongoDB, only ops and infrastructure cost
- Very little to zero data loss in failure
- Down for a very short interval
- BUT... programmer errors will replicate almost instantly

Application Engineering

- Flexible
- Expensive, time and money
- Good use of engineer?

Backup

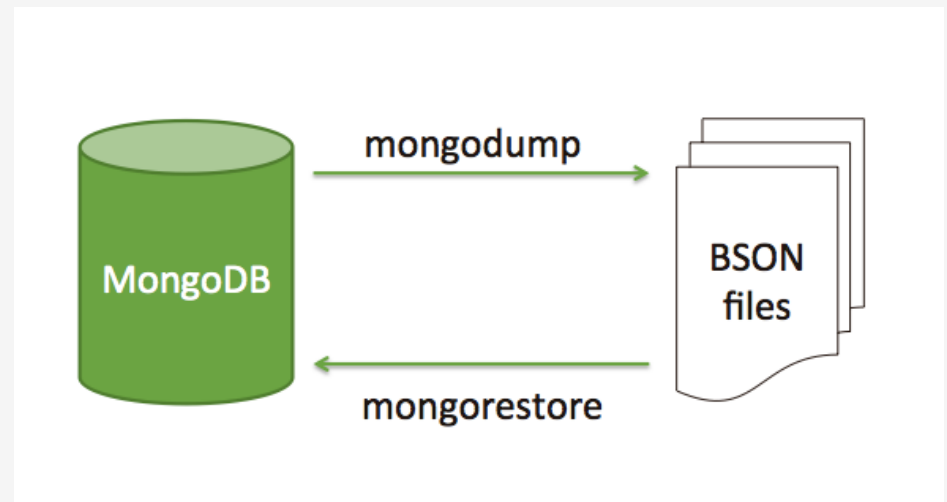
- Slower to restore
- Can suffer from being out of date
- Fairly cheap
- Isolated
- Covers most risks

Backup Approaches

- Mongodump/mongorestore
- Storage level options
- MongoDB Backup Service

Mongodump/MongoRestore

- Can be run online or offline
- Oplog aware for point in time restores
- Filter in, filter out
- Considerations
 - Data size
 - Sharding



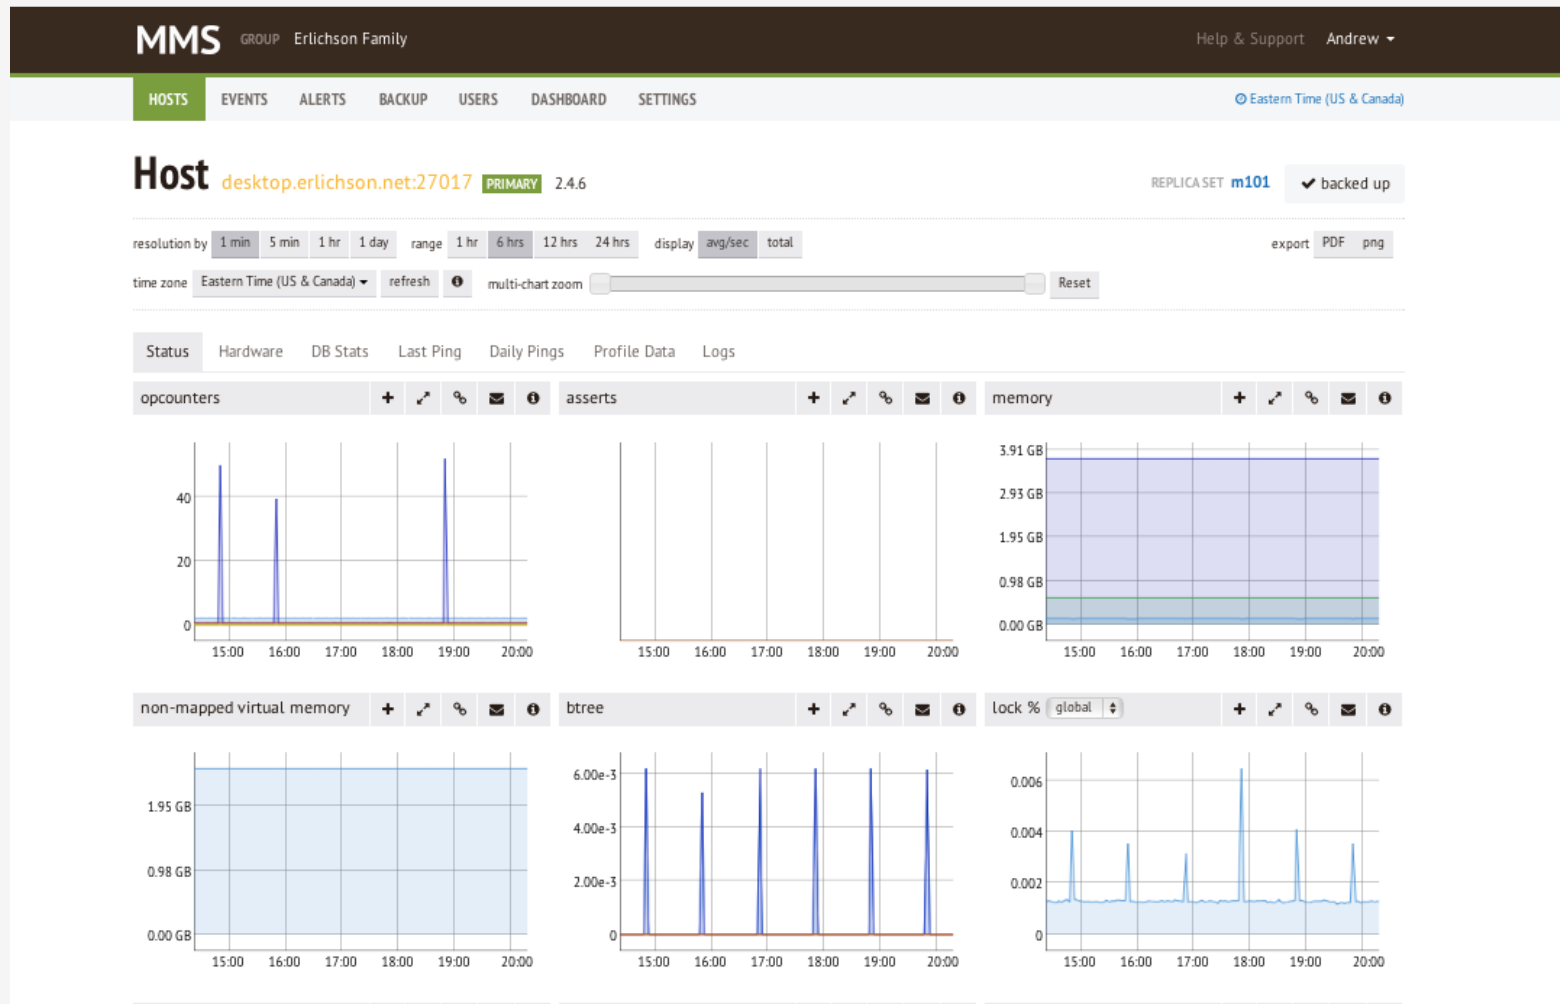
Storage-level Backups

- Copy files in your data directory (e.g. /data/db)
- File system or block storage snapshots
- Fastest way to backup/restore
- Considerations
 - Journal
 - Consistency
 - Backup granularity (whole file system back up?)
 - Ops expertise
 - Storage of snapshots or data file backups

What is MongoDB Management Service (MMS)?



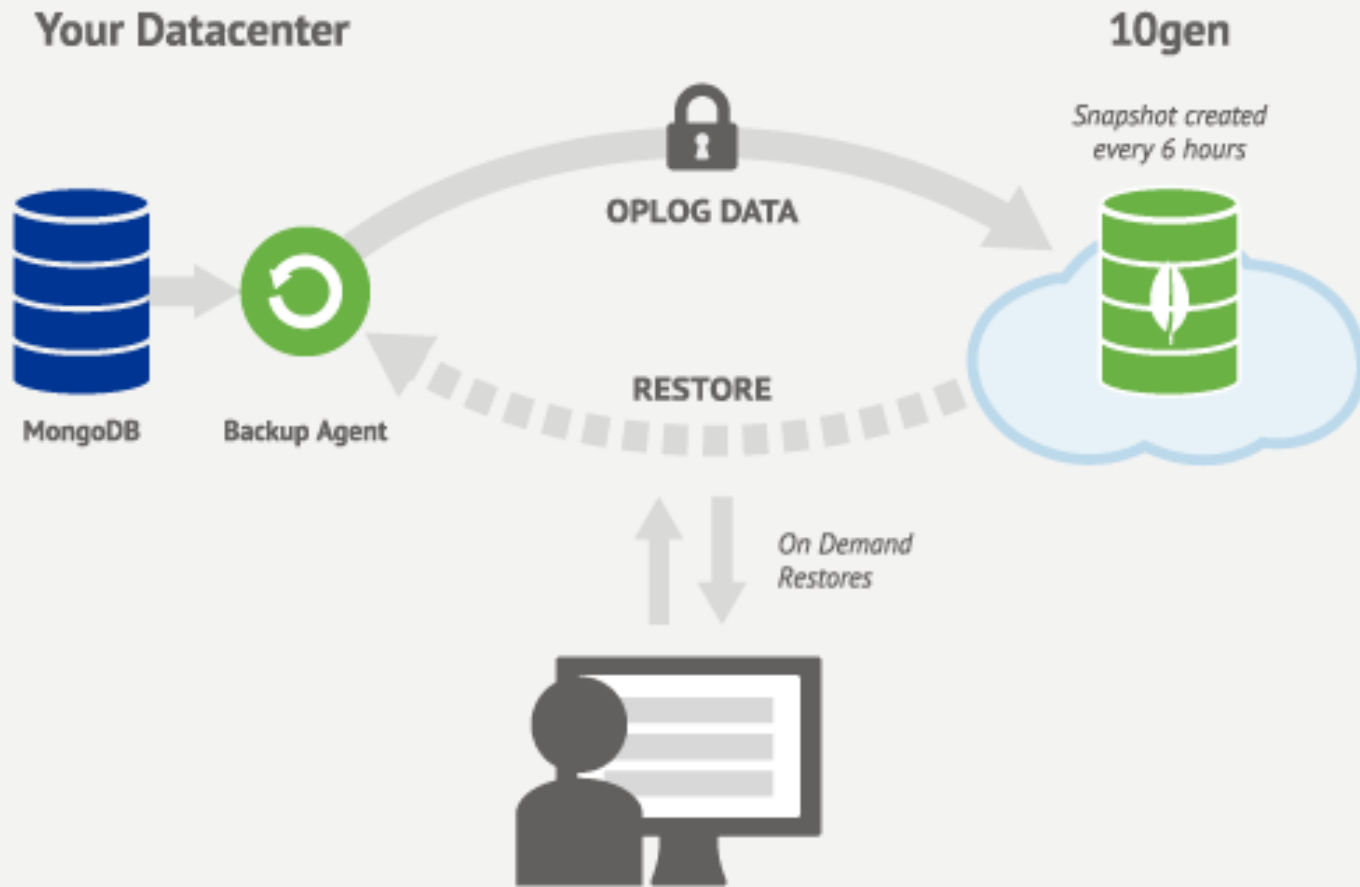
Free monitoring



How it Works



Overview



Under the Hood

- From the initial sync, we rebuild your data in our datacenters and take a snapshot
- We take snapshots every 6 hours
- Oplog is stored for 48 hours

Sharded Clusters

- Balancer paused every 6 hours
- A no-op token is inserted across all shards, mongoses and config servers
- Oplog applied to replica sets until point in which token was inserted
- Provides a consistent state of database across shards

Recovery Approaches

Recovery Approaches

	Mongodump	File system	MMS Backup
Initial complexity	Medium	High	Low
Confidence in Backups	Medium	Medium	High
Point in time recovery of replica set	Sort of 😊	No	Yes
System Overhead	High	Can be low	Low
Scalable	No	With work	Yes
Consistent Snapshot of Sharded System	Difficult	Difficult	Yes

Key Benefits



Point in Time

- Oplog stored for 48 hours
- Restore your replica set to any point-in-time in the last 48 hours by creating a custom snapshot

Snapshot Policy

- Every 6 hours for 48 hours
- Every day for 1 week
- Every week for for a month
- Every month for a year

Easy to Restore

- Pull from custom URL
- Push via SCP

Unlimited Restores

- Confidence in your restore process
- Build development, QA, analytics environments without impacting production

Fully Managed

- Created by the engineers that build MongoDB
- No need to write or maintain custom backup scripts

Getting Started



Getting Started

1. Create an account at mms.mongodb.com
2. Install MMS Monitoring on your deployment
3. Register at mms.mongodb.com/backup
4. Install MMS Backup on your deployment
5. Start initial sync
6. Rest easy!

**Free Month – Register Using
MONGODB_1MONTHFREE**

Questions?





#ParisMongoDB

Thank You

Sam Weaver

Solution Architect, MongoDB