

# { Ops : Workshop }

MongoDB London  
Sam Weaver  
@samuel\_weaver



# Agenda

- Intro to MongoDB
- Replica Sets
- Sharding
- Other topics
- Hands on!
- Questions? Ask!



# Getting Started

- Download mongod
- Windows, Linux, Mac OSX
- <http://www.mongodb.org/downloads>
- Odd and even version numbers?
- Unzip and go!



# Running a single mongod

- Make a directory for your data

```
mkdir -p /data/db
```

- Start up mongod

```
mongod --dbpath /data/db
```

- --logpath, --logappend, --fork, --rest, --port



# Connect and use

- `mongo`
- `show dbs`
- `show collections`
- `use <database>`
- `db.test.insert()`
- `db.test.find()`



# Database files

- Check /data/db
- test.ns 16MB
- test.0 64MB
- test.1 128MB
- “Huh? What’s up with all these files?”



# Exercise

- Download and install mongod
- Connect to the shell
- Create a new “workshop” database
- Create a document in the “test” collection



# Replica Sets

- What are replica sets?
  - Quorum, heartbeat
- Primary and secondary's
  - Replication
  - Oplog & oplog sizing
- Automatic Failover
  - Voting
- Write concern





# Setting up a replica set

- make a data directory for each node

```
mongod --dbpath /data/node1 --replSet test --port 27017
```

```
mongod --dbpath /data/node2 --replSet test --port 27018
```

```
mongod --dbpath /data/node3 --replSet test --port 27019
```

- `rs.initiate()`
- `rs.status()`
- `rs.add()`



# Config

- `conf = rs.conf()`
- `conf.members[2].priority = 100`
- `rs.reconfig(conf)`
- Options
  - Hidden
  - Slave delay
  - Priority
  - Arbiter



# Tagging

```
{
  _id : "mySet",
  members : [
    { _id : 0, host : "A", tags : { "dc" : "ny" } },
    { _id : 1, host : "B", tags : { "dc" : "ny" } },
    { _id : 2, host : "C", tags : { "dc" : "sf" } },
    { _id : 3, host : "D", tags : { "dc" : "sf" } },
    { _id : 4, host : "E", tags : { "dc" : "cloud" } } ],
  settings : {
    getLastErrorModes : {
      allDCs : { "dc" : 3 },
      someDCs : { "dc" : 2 } }
  }
}
> db.blogs.insert({...})
> db.runCommand({getLastError : 1, w : "someDCs"})
```

# Other commands

- `rs.help()`
- `rs.status()`
- `rs.slaveOk()`
- `db.printReplicationInfo()`
- `db.printSlaveReplicationInfo()`

# Exercises

- Set up a 3 node replica set
- Run the command to step down a primary: `rs.stepDown()` and ensure that a secondary is elected the new primary. Bring the set back up to 3 members again
- Set a priority on a node. Terminate the primary node and practice automated failover, see what happens.
- Insert some data into primary and read it from a secondary - what happens?
- Add a new node (so 4 members now) in your RS and kill 2, does a primary get elected? See what happens.



# Sharding

- What is sharding?
- Why shard?
- Sharded architecture
- Chunks
- MongoS balancing



# Sharding setup

- **make a directory for each shard and a config database**

```
mkdir /data/shard1 /data/shard2 /data/config
```

- **start a config server**

```
mongod --configsvr
```

- **start a mongos**

```
mongos --configdb "xxxx.local" --chunkSize 1
```

- **start 2 mongod's**

```
mongod --dbpath /data/shard1 --port 27020
```

```
mongod --dbpath /data/shard2 --port 27021
```



# Connecting to a shard

- mongo
- Will prompt with mongos
- `sh.status()`
- `sh.addShard()`
- `sh.enableSharding()`
- `sh.shardCollection()`





# Exercises

- Setup a 2 shard cluster
- Insert some test data and shard on a particular field
- Watch what happens with the balancing
- Explore the configdb



# Backup and Restore

- Mongodump/mongorestore
- Mongoexport/mongoimport
- File system snapshot (LVM)
- File copy



# Exercises

- Use mongodump to dump your test database from a single mongod
- Use mongorestore to restore it
- Use mongoexport to export your test database to json
- Use mongoimport to import your database
- Mongodump from a sharded set up



# Monitoring Server

- MMS
- Munin/Nagios
- iostat -xm 2
- top
- mongostat



# Monitoring MongoDB

- `db.serverStatus()`
- `db.stats()`
- `db.test.stats()`
- `db.currentOp()`



# Exercise

- Create a for-loop to insert data into MongoDB

```
for(x=0;x<10000<x++) { db.test.insert({x:x}) }
```

- Watch mongostat statistics
- Find out the current operation of MongoDB
- Find out how much memory and cpu is being utilized



# Production Notes

- 64-bit operating system
- ext4 or efs file system
- focus on RAM not cpu
- disk type (SAS 15k, SSD) for best performance
- no atime
- ulimit = 20,000
- No NUMA, no hugepages
- 16kb readahead
- RAID 10



# Thank you

@samuel\_weaver

